

# CITS5501 Software Testing and Quality Assurance Alloy

Unit coordinator: Arran Stewart

- More on motivation for formal methods
- Alloy specification language and Alloy analyser

## Issues with specifications

System specifications can suffer from a few potential problems:

- contradictions
- ambiguities
- vagueness
- incompleteness

(NB: these terms all have precise meanings. Don't assume they're interchangeable.)

1. [Introduction](#)

- *Contradictions.* In a very large set of specifications, it can be difficult to tell whether there are requirements that contradict each other.
  - Can arise where e.g. specifications are obtained from multiple users/stakeholders
  - Example: one requirement says “all temperatures” in a chemical reactor must be monitored, another (obtained from another member of staff) says only temperatures in a specific range.

## Some problems in specifying systems

- *Ambiguities.* i.e., statements which can be interpreted in multiple different ways.  
“The operator identity consists of the operator name and password; the password consists of six digits. It should be displayed on the security screen and deposited in the login file when an operator logs into the system.”

## Some problems in specifying systems

- *Ambiguities.* i.e., statements which can be interpreted in multiple different ways.  
“The operator identity consists of the operator name and password; the password consists of six digits. It should be displayed on the security screen and deposited in the login file when an operator logs into the system.”
- ... Does “it” refer to the identity, or the password?

## Some problems in specifying systems

- *Ambiguities.* i.e., statements which can be interpreted in multiple different ways.

“The operator identity consists of the operator name and password; the password consists of six digits. It should be displayed on the security screen and deposited in the login file when an operator logs into the system.”
- ... Does “it” refer to the identity, or the password?
- “Should” can be ambiguous – does “The system should do *X*” mean the system *must* do *X*, or that it is optional but desirable that the system do *X*?



## Some problems in specifying systems

- *Ambiguities*. i.e., statements which can be interpreted in multiple different ways.

“The operator identity consists of the operator name and password; the password consists of six digits. It should be displayed on the security screen and deposited in the login file when an operator logs into the system.”
- ... Does “it” refer to the identity, or the password?
- “Should” can be ambiguous – does “The system should do *X*” mean the system *must* do *X*, or that it is optional but desirable that the system do *X*?
- Many terms have both technical and non-technical meanings (possibly multiple of each): for instance, “reliable”, “robust”, “composable”, “category”, “failure”, “orthogonal”, “back end”, “kernel”, “platform”, “entropy” ...

## Some problems in specifying systems

- *Vagueness*. Vagueness occurs when it's unclear what a concept covers, or which things belong to a category and which don't.

- *Vagueness*. Vagueness occurs when it's unclear what a concept covers, or which things belong to a category and which don't.
- "is tall" is vague: some people are definitely tall, and some are definitely short, but it can be difficult to tell when exactly someone meets the criterion of being tall.

## Some problems in specifying systems

- *Vagueness*. Vagueness occurs when it's unclear what a concept covers, or which things belong to a category and which don't.
- "is tall" is vague: some people are definitely tall, and some are definitely short, but it can be difficult to tell when exactly someone meets the criterion of being tall.
- Likewise "fast", "performant", "efficiently", "scalable", "flexible", "is user-friendly", "should be secure", "straightforward to understand" are all vague.

## Some problems in specifying systems

- *Incompleteness*. This covers specifications that, for instance, fail to specify what should happen in some case.

## Some problems in specifying systems

- *Incompleteness*. This covers specifications that, for instance, fail to specify what should happen in some case.
- e.g. An obviously incomplete requirement: “A user may specify normal or emergency mode when requesting a system shutdown. In normal mode, pending operations shall be logged and the system shut down within 2 minutes.”

## Some problems in specifying systems

- *Incompleteness*. This covers specifications that, for instance, fail to specify what should happen in some case.
- e.g. An obviously incomplete requirement: “A user may specify normal or emergency mode when requesting a system shutdown. In normal mode, pending operations shall be logged and the system shut down within 2 minutes.”
- ... So what happens in emergency mode?

## Some problems in specifying systems

- *Incompleteness*. This covers specifications that, for instance, fail to specify what should happen in some case.
- e.g. An obviously incomplete requirement: “A user may specify normal or emergency mode when requesting a system shutdown. In normal mode, pending operations shall be logged and the system shut down within 2 minutes.”
- ... So what happens in emergency mode?
- But other cases of incompleteness may be harder to spot.



## Some problems in specifying systems

- In addition to these, there are many other ways requirements can be written poorly –  
e.g. Overly long and complex sentences, mixed levels of abstraction (mixing high-level, abstract statements with very low-level ones → difficult to distinguish high-level architecture from low-level details), undefined jargon terms, specifying implementation rather than requirements (how vs what), over-specifying, don't satisfy business needs, etc.

## Some problems in specifying systems

- In addition to these, there are many other ways requirements can be written poorly –  
e.g. Overly long and complex sentences, mixed levels of abstraction (mixing high-level, abstract statements with very low-level ones → difficult to distinguish high-level architecture from low-level details), undefined jargon terms, specifying implementation rather than requirements (how vs what), over-specifying, don't satisfy business needs, etc.
- Formal specifications can potentially help avoid ambiguity, vagueness, contradiction and some gaps in completeness.

## Some problems in specifying systems

- In addition to these, there are many other ways requirements can be written poorly –  
e.g. Overly long and complex sentences, mixed levels of abstraction (mixing high-level, abstract statements with very low-level ones → difficult to distinguish high-level architecture from low-level details), undefined jargon terms, specifying implementation rather than requirements (how vs what), over-specifying, don't satisfy business needs, etc.
- Formal specifications can potentially help avoid ambiguity, vagueness, contradiction and some gaps in completeness.
- Other problems, not so much. Just as it's possible to write programs badly in any language, it's also possible to write formal specifications badly.

There is still a need for *review* of specifications, as with any artifact.

## Some problems in specifying systems

Some example requirements<sup>1</sup>:

- “The system shall be highly reliable.”
- “Credit card details must be encrypted.”
- “The system should be easy to use by medical staff and should be organized in such a way that user errors are minimized.”
- “The system shall have a response time of under 0.1 seconds.”
- “If the LAUNCH-MISSILE signal is set to TRUE and the ABORT-MISSILE signal is set to TRUE then do not launch the missile, unless the ABORT-MISSILE signal is set to FALSE and the ABORT-MISSILE-OVERRIDE is also set to FALSE, in which case the missile is not to be launched.”

<sup>1</sup>Sources: Pressman; LaPlante, *Requirements Engineering*; bitter experience.

- Formal specifications can help with ameliorating these problems.
- Sometimes, just the process of attempting to formalize a requirement can reveal problems with it.
- Using a formal model can help reveal *ambiguity* and *vagueness* and allow them to be eliminated
- It may also be possible (depending on the mathematical model used) to detect inconsistencies
- Detecting whether a specification is *complete* is more difficult.
  - Some gaps may be able to be detected
  - But there are nearly always some details that are left undefined, or scenarios that may not have been considered.

## Formal specifications

- Formal specifications have a meaning defined in terms of mathematics.
- Similar to programming languages, there are different sorts of formal specification languages and tools, with different sorts of scope.
- Some are small and specific in scope. For instance:
  - *State charts*, which we have seen, define states and transitions between them.
  - *Backus-Naur Form*, which we have also seen, defines languages – sets of strings.
  - *Regular expressions* also define sets of strings (most programming languages have an implementation of them)
  - $\pi$ -calculus is used to represent concurrent systems
- Highlights another reason to use formal specifications: sometimes they're the most concise way of saying something. Writing a grammar or set of state transitions in natural language (rather than BNF or as a state chart) would be a gruelling and horrible exercise.
- Some are more general, and are intended to be able to describe a wide range of systems. They may be based on numerous different mathematical formalisms – predicate logic + set theory, category theory, etc.

# Formal specification languages

Some examples of general-purpose specification languages:

- **Z notation**

- based on set theory and predicate logic
- developed in the 1970s.
- Now has an ISO standard, and variations (e.g. object-oriented versions)

- **TLA+:**

- Stands for “Temporal Logic of Actions”
- A general-purpose specification language
- Especially well-suited for writing specifications of concurrent and distributed systems
- For finite state systems, can check (up to some number of steps) that particular properties hold (e.g. safety, no deadlock)

100%

- We'll be using the **Alloy** specification language
- Alloy is both a language for describing structures, and a tool (written in Java) for exploring and checking those structures.
- Influenced by Z notation, and modelling languages such as UML (the Unified Modelling Language).
- Website: <http://alloy.mit.edu/> (The Alloy Analyzer tool can be downloaded from here.)



## Alloy language and analyser

The idea behind Alloy is that:

- It lets you capture the essence of a design at a high level
- It lets you identify risky aspects of a design
- It lets you develop a model incrementally
- It lets you simulate and analyze the model as you go

# Alloy idea

In other words, *before* you start implementing a system, you can start specifying the entities that make it up, what constraints (i.e. invariants) hold for them, and how they hang together.

# What is Alloy?

- A flexible *language* for describing structures (and how they interrelate)
- It can describe both
  - static structures
  - dynamic behaviours<sup>2</sup>

<sup>2</sup>Alloy now also has an extension, Electrum, for modelling properties of systems over time using temporal logic, but we will restrict ourselves to plain Alloy.

# What is Alloy?

- A flexible *language* for describing structures (and how they interrelate)
- It can describe both
  - static structures
  - dynamic behaviours<sup>2</sup>
- Comes with a tool, the Alloy *Analyzer*
  - Generates counterexamples to theorems/statements

---

<sup>2</sup>Alloy now also has an extension, Electrum, for modelling properties of systems over time using temporal logic, but we will restrict ourselves to plain Alloy.

100%

- Small and easy to use
- Has a simple and uniform semantics based on mathematical *relations*
- Can be easily analysed using automated tools

# Comparison with UML

Alloy has some similarities with UML –

- It has a graphical notation
- It is somewhat similar to the *Objects Constraint Language* use by UML<sup>3</sup>

And several differences:

- Unlike UML, Alloy has precise semantics
- It is a far smaller and simpler formalism than UML
  - UML allows for many constructs (e.g. use cases, state charts) that don't have an equivalent in Alloy

---

<sup>3</sup>[https://en.wikipedia.org/wiki/Object\\_Constraint\\_Language](https://en.wikipedia.org/wiki/Object_Constraint_Language)

## Using Alloy

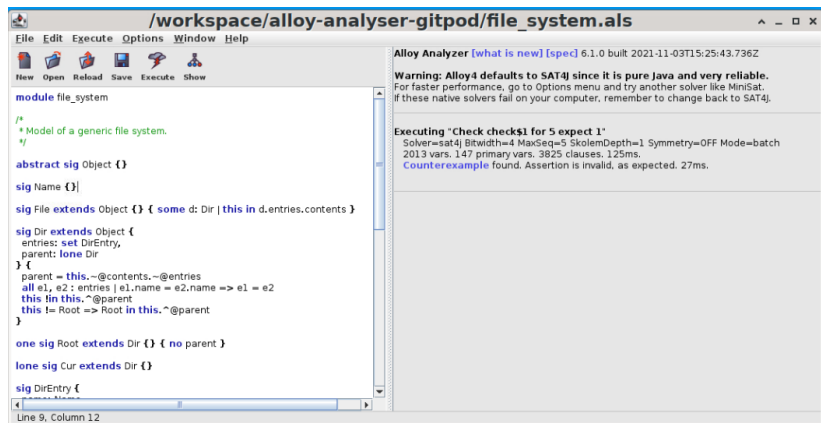
- The Alloy analyser is distributed as a Java `.jar` file (or a `.dmg` file for Mac OS X) – see the [Alloy 6.0 release page](#)
  - The `.jar` file can be run like this:

```
java -jar org.alloytools.alloy.dist.jar
```



# Using Alloy

I have also set up a GitHub repository which lets you use the Alloy analyser from within an online IDE using [Gitpod](https://github.com/arranstewart-dev/alloy-analyser-gitpod) – visit <https://github.com/arranstewart-dev/alloy-analyser-gitpod/>



The screenshot shows the Alloy Analyzer web IDE interface. The title bar indicates the workspace path: `/workspace/alloy-analyser-gitpod/file_system.als`. The menu bar includes File, Edit, Execute, Options, Window, and Help. The toolbar contains icons for New, Open, Reload, Save, Execute, and Show. The main editor displays the following Alloy code:

```
module file_system

/*
 * Model of a generic file system.
 */

abstract sig Object {}

sig Name {}

sig File extends Object {} {
  some d: Dir | this in d.entries.contents
}

sig Dir extends Object {
  entries: set DirEntry,
  parent: lone Dir
} {
  parent = this.~@contents.~@entries
  all e1, e2: entries | e1.name = e2.name ==> e1 = e2
  this in this.~@parent
  this != Root ==> Root in this.~@parent
}

one sig Root extends Dir {} { no parent }

lone sig Cur extends Dir {}

sig DirEntry {
  name: Name
}
```

The right-hand pane shows the Alloy Analyzer version (6.1.0) and build information. It displays a warning that Alloy4 defaults to SAT4j and a message indicating that the analysis was successful, finding a counterexample for the assertion.

Alloy Analyzer [what is new] [spec] 6.1.0 built 2021-11-03T15:25:43.736Z

Warning: Alloy4 defaults to SAT4j since it is pure java and very reliable. For faster performance, go to Options menu and try another solver like MiniSat. If these native solvers fail on your computer, remember to change back to SAT4j.


Executing "Check check\$1 for 5 expect 1"

Solver=sat4j Bitwidth=4 MaxSeq=5 SkolemDepth=1 Symmetry=OFF Mode=batch 2013 vars. 147 primary vars. 3825 clauses. 125ms.

Counterexample found. Assertion is invalid, as expected. 27ms.

Line 9, Column 12

## Using Alloy



## Alloy analyser displaying a counterexample

## Alloy

In Alloy, we declare rules about a mini-universe: things that exist, and properties that should be true of them.

- “There are things called animals”

```
sig Animal {}
```

- “A cat is a sort of animal”

```
sig Cat extends Animal {}
```

## Alloy

When modelling entities in Alloy – we normally include only the bare minimum of properties needed in order to show how the system “hangs together”.

# Alloy language

- For example – we'll look at a simple model of a file system (based on the Alloy tutorial at <http://alloytools.org/tutorials/online/>)
- The Alloy specification looks a little like Java:

```
// A file system object in the file system
sig FSObject { parent: lone Dir }
```

```
// A directory in the file system
sig Dir extends FSObject { contents: set FSObject }
```

```
// A file in the file system
sig File extends FSObject { }
```

# Alloy primitives

- In Alloy, everything is built up from **atoms** and **relations**
- An *atom* in alloy is an indivisible, immutable value
  - We don't create these directly – they get automatically generated by the analyser
  - Example atoms: A0, A1, B0, R0 ...
- A *relation* is a structure that relates atoms together –
  - It is a set of **tuples**

## Alloy primitives

The easiest way to think of relations is probably to think of them as a sort of table – which show how columns of things are **related**.

e.g. “shares an office with”:

Person A	Person B
Alice	Bob
Bob	Alice
Dan	Eve
Eve	Dan

Each row is called a *tuple*.

**CNN**



# Alloy – relations

Alloy's semantics are defined in terms of relations.

Example relations:

- “Is less than”. e.g. “ $2 < 4$ ”, “ $10 < 9$ ”.
- “Is the blood relative of”. e.g. “Alice is the blood relative of Bob”.
- “Shares an office with”. e.g. “Bob shares an office with Carol”.

These are all *binary relations*. Statements about two entities, which can be true or false.

# Alloy – relations

Relations can also be *unary* (about one entity):

- “Is even”. e.g. “*even*(2).”
- “Is an employee”. e.g. “Dan is an employee”.

# Alloy – relations

They can be ternary:

- “\_ is delivered to \_, by \_”. e.g. “The *blue book* was delivered to *Alice*, by *Bob*”.

Or, in general, they can be  $n$ -ary – a statement about  $n$  things.

# Alloy – relations

We can think of predicates as being functions – an  $n$ -ary predicate isn't true or false in itself, until we supply it with  $n$  arguments.

- “Is less than” isn’t true or false, but “ $2 < 4$ ” is.

# Alloy – relations

Relations can be finite, or infinite.

An infinite relation: “is less than”

Number A	Number B
1	2
1	3
2	3
...	...

# Alloy relations

- *Sets* are unary (1-column) relations. e.g.

```
Name = { N0,  
          N1,  
          N2 }
```

- *Scalars* are actually 1-element sets:

```
myName = N0
```

- Binary or ternary or higher relations are possible:

```
names = { (B0, N0),  
          (B0, N1),  
          (B1, N2) }
```

# Alloy – sigs

`sig Animal {}` says “There are things called animals”.

It defines a unary relation, “Animal”. Something thing can be-an-animal, or not.

# Alloy – sigs

`sig Cat extends Animal {}` says “Cats are a sort of animal”.

*If* something has the property “is-an-animal”, *then* it might also have the property “is-a-cat”.

We can read “extends” as also meaning “is a kind of”, or “is a subtype of”.



# Alloy – subtypes

- So, **extends** indicates subtypes (similar to Java).
- Here, **Dir** and **File** are both subtypes of **FSObject**:

```
sig FSObject {}
```

```
sig Dir extends FSObject {}
```

```
sig File extends FSObject {}
```

- When we declare **Dir** or a **File** to be sub-types of **FSObject**, they are considered to be *mutually disjoint* sets
- The above says “There are things called FSObjects. An FSObject might be a Dir or it might be a File, but not both”.

# Alloy – properties

We can specify *properties* of entities:

```
// A file system object in the file system  
sig FSObject { parent: lone Dir }
```

```
// A directory in the file system  
sig Dir extends FSObject { contents: set FSObject }
```

```
// A file in the file system  
sig File extends FSObject { }
```

# Alloy – properties

```
// A file system object in the file system  
sig FSObject { parent: lone Dir }
```

```
// A directory in the file system  
sig Dir extends FSObject { contents: set FSObject }
```

```
// A file in the file system  
sig File extends FSObject { }
```

These are usually written within the sig of an entity.

They actually represent *relations* between entities.

# Alloy – properties

```
// A file system object in the file system  
sig FSOobject { parent: lone Dir }
```

There are multiple ways of reading this:

- “There are such things as FSOobjects. An FSOobject has the property ‘parent’. An FSOobject can have zero or one parents.”  
Or –
- “A relation ‘parent’ exists between FSOobjects and Dirs. Whenever an FSOobject appears in the relation, it can be association with at most one Dir.”

These are exactly equivalent.

## Alloy – properties

```
// A file system object in the file system  
sig FSObject { parent: lone Dir }
```

- The “lone” means “zero or one”. It is a *cardinality*.
- Other possible cardinalities are:
  - “some” (one or more)
  - “one” (exactly one)
  - “set” (zero or more)
- When we specify a property using a colon in this way, the default multiplicity is one.
- We can use cardinalities whenever we are specifying a set or relation: since sigs also represent sets (e.g. the set of Dirs), we can give them cardinalities, too.

## Alloy – properties

```
one sig RootDir extends Dir { }
```

There exists a “RootDir”, but only one of them.

# Exercise

Games:

- There are things called games.
- Games can be board games, or field games.
- There may be other sorts of games.

# Alloy language – comments

```
// A file system object in the file system  
sig FSObject { parent: lone Dir }
```

```
// A directory in the file system  
sig Dir extends FSObject { contents: set FSObject }
```

```
// A file in the file system  
sig File extends FSObject { }
```

- Comments can be written in multiple ways



# Alloy language – comments

```
// A file system object in the file system  
sig FSObject { parent: lone Dir }
```

```
// A directory in the file system  
sig Dir extends FSObject { contents: set FSObject }
```

```
// A file in the file system  
sig File extends FSObject { }
```

- Comments can be written in multiple ways
  - single-line comments with “//” or “- -”

# Alloy language – comments

```
// A file system object in the file system  
sig FSObject { parent: lone Dir }
```

```
// A directory in the file system  
sig Dir extends FSObject { contents: set FSObject }
```

```
// A file in the file system  
sig File extends FSObject { }
```

- Comments can be written in multiple ways
  - single-line comments with “//” or “- -”
  - multiline comments with “/\* ... \*/”

# Alloy – facts

- How can we express that any **FSObject** is either a **Dir** or a **File**?  
(i.e., there are no other sorts of **FSObject**)

# Alloy – facts

- How can we express that any **FSObject** is either a **Dir** or a **File**? (i.e., there are no other sorts of **FSObject**)
- Alloy also allows us to specify *constraints*. These are introduced with the keyword **fact**.

```
sig FSObject { parent: lone Dir }  
sig Dir extends FSObject { contents: set FSObject }  
sig File extends FSObject { }
```

```
// All file system objects are either files or directories  
fact { File + Dir = FSObject }
```

- The general syntax for a fact is  
**fact** *name* { *formulas* }
- *formulas* are Boolean expressions, and by putting them in a fact, we're constraining them to be true.

## Alloy – abstract signatures

- (An alternative way to say that all FSObjects must be Dirs or Files would be to declare FSObject **abstract**)

## Alloy – abstract signatures

- (An alternative way to say that all FSObjects must be Dirs or Files would be to declare FSObject **abstract**)
- (This is similar to the use of the **abstract** keyword in Java; it means there are no objects that are *directly* of type FSObject; they must be members of some subtype, instead.)

## Alloy – operators

Operators are available to construct Boolean expressions.

- subset: **in**
  - *set1 in set2* — *set1* is a subset of *set2*
  - informally: “some *set2* are *set1*”, or “a *set2* may be *set1*”; but the set-theoretic meaning is more precise.
- set equality: **=**
  - *set1 = set2* — *set1* equals *set2*
- scalar equality: **=**
  - *scalar = value* — *scalar* equals *value*



## Alloy – subsets

- We saw that subtypes are disjoint.
- We can also declare *subsets*:

```
sig signame in supername { ... }
```

- Subsets are *not* necessarily disjoint, and may have multiple parents





# Relations

- ```
// A file in the file system
sig File extends FSObject { }
```

# Relations

- ```
// A file system object in the file system
sig FSObject { parent: lone Dir }
```
- ```
// A directory in the file system
sig Dir extends FSObject { contents: set FSObject }
```
- ```
// A file in the file system
sig File extends FSObject { }
```
- To a first approximation, we can think of relations as behaving like *fields* in an OO language.
- `sig FSObject { parent: lone Dir }` can be read as “Things of type `FSObject` *have a parent*, which is of type `Dir`”.
- **lone** means “at most one” – i.e., you can have zero or one parents. (We need this because the root directory has no parent.)

# Relations

- ```
// A file system object in the file system
sig FSObject { parent: lone Dir }

// A directory in the file system
sig Dir extends FSObject { contents: set FSObject }

// A file in the file system
sig File extends FSObject { }
```
- More precisely, parent is a relation between FSObject and Dir.

## Relations – multiplicities

- **lone** is a type of multiplicity – it says how many of something there are.
- Other multiplicities:
  - **one** - one
  - **some** - at least one; one or more
  - **set** - zero or more
  - **no** - zero
- The default multiplicity is one.

## Relations – multiplicities

- In set theory terms ...
  - **one** means the relation is a total function –  
sig Student { name : one String } –  
for every Student, we can map to a string which is their name.
  - **lone** means the relation is a *partial* function –  
sig Student { driverLicenseNum : lone String } – \
- for every Student, we *may* be able to map to a diver's license number.
- (Here, it's assumed you can't have more than one license.)



# Relations

- So, signature declarations will look like:

```
sig SomeName {
  field1 : FieldType,
  field2a, field2b : OtherFieldType
}
```

- The order of declarations doesn't matter – `SomeName`, `FieldType` and `OtherFieldType` could be declared in any order in a file.

# Relations

- `// A directory in the file system`  
`sig Dir extends FSObject { contents: set FSObject }`
- Here, we say that a `Dir` has a field `contents`, which is a *set* of `FSObjects`.
- The could contain one item, many items, or no items.

## Examples

- “A car has one engine”  

```
sig Car { engine: one Engine }, or  
sig Car { engine: Engine }
```
- “People have zero or more hobbies”  

```
sig Person { hobbies: set Activity }
```

## Exercises

- Classes have at least one lecturer, and zero or more students.

## Exercises

- Classes have at least one lecturer, and zero or more students.
- Animals have zero or more legs

## Exercises

- Classes have at least one lecturer, and zero or more students.
- Animals have zero or more legs
- Some animals are carnivores

## Exercises

- Classes have at least one lecturer, and zero or more students.
- Animals have zero or more legs
- Some animals are carnivores
- Textbooks have one or more pages


---

- FSObjects have parents, and directories have contents, and we have constrained the multiplicities ...
- but there's currently no connection between them.



# File system

- So we could have this situation:



# File system

- We will need to constrain things more, so we'll use a *fact*.

```
// A directory is the parent of its contents
fact { all d: Dir, o: d.contents | o.parent = d }
```

- This says: “for any thing (let’s call it  $d$  for the moment) of type `Dir`, and for any thing (let’s call it  $o$  for the moment) which is in the set `d.contents`:  
 $o$ ’s parent is  $d$ .”
- It uses a *quantifier* (“all”) – we’ll look at these more later.